

Yash Kherwal

+91 9891148156 | yash.kherwal78@gmail.com | linkedin.com/in/yash-kherwal-944497254 | github.com/YashKherwal78 |
careerautomated.in

EDUCATION

Indian Institute of Technology Roorkee

2022 – 2026

B.Tech in Chemical Engineering

EXPERIENCE

AI Engineer Intern

Feb 2026 – Apr 2026

Orange Labs — EdTech ERP platform for schools and colleges

- Developed computer vision model detecting students from live camera feeds, auto-syncing presence records into the school ERP, eliminating manual marking.
- Engineered synthesised audio/video from real teacher input, improving engagement over generic video content.

Software Development Intern

May 2025 – June 2025

ScoreMe Solutions

- Implemented **confidence calibration layer**: predictions below threshold withheld from **credit scoring pipelines** and flagged for review; integrated **Tesseract OCR** as fallback for scanned inputs.
- Designed rule-based heuristics for fintech PDF processing, automating processing and avoiding unnecessary model inference on clear-cut documents.

Engineering Intern

June 2025 – July 2025

Bharat Electronics Limited

- Implemented a **binary stream parser** for **ASTERIX CAT048** radar protocol per EUROCONTROL spec; designed **concurrent ingestion** with **backpressure handling** to maintain throughput under sustained high-frequency surveillance loads.

PROJECTS

CareerAutomated — Autonomous AI Recruiting Platform | *Python, Playwright, LangGraph, RAG*

2025 – 2026

- Shipped an **LLM-based resume tailoring engine** under an **integrity gate** that fact-checks claims against source resumes
- Developed an end-to-end AI job platform: **50+ ATS discovery connectors** ingesting and normalizing **1.3M+ live job postings** with a **hybrid BM25 + semantic RAG** pipeline
- Engineered a **plugin-based auto-apply engine** with a shared Playwright core and per-ATS handlers

Hybrid RAG Retrieval System | *Python, LangChain, BGE-M3, AstraDB, FastAPI, AWS EC2*

2024

- Diagnosed keyword retrieval failure on 500+ documents and introduced **hybrid retrieval** combining **BGE-M3** with BM25 scores
- Designed **decoupled ingestion and serving** on **AWS EC2** for zero-downtime corpus updates with **response groundedness** constraints

YAAR — Personality-Adaptive LLM System | *React Native, LLM APIs, FastAPI, Persistent Memory*

2025

- Shipped production **Android APK** via React Native + EAS Build — validated end-to-end on device with sub-2s streamed LLM response latency.
- Designed a **persistent personality memory system** updating a **personality metric vector** injected into the **LLM context window** at inference time

TECHNICAL SKILLS

LLM & Agentic AI: LangGraph, LangChain, CrewAI, Multi-Agent Orchestration, MCP, Structured Outputs, Prompt Engineering, Agentic Workflows

Retrieval: Hybrid RAG, BGE-M3, Dense + Sparse Retrieval, AstraDB, Vector Databases, ANN Indexing

ML / AI: TensorFlow, Scikit-learn, Random Forest, Confidence Calibration, Computer Vision (OpenCV)

Backend & Infra: FastAPI, Docker, AWS EC2, n8n, REST APIs, Stream Processing

Languages & Data: Python, SQL, Java, C++, DSA, Pandas, NumPy