

Yash Kherwal

+91 9891148156 | yash.kherwal78@gmail.com | linkedin.com/in/yash-kherwal-944497254

EDUCATION

Indian Institute of Technology Roorkee

2022 – 2026

B.Tech in Chemical Engineering

EXPERIENCE

AI Product Manager Intern

Feb 2026 – Apr 2026

OrangeLabs (EdTech)

- Owned end-to-end product development of two AI features — defined the problem, designed the solution, and worked directly with engineering to ship; **AI Attendance System** (CCTV-based auto-detection) eliminated manual roll calls, freeing ~10 minutes of instruction time per class period.
- Conducted **customer discovery** with school administrators and teachers; identified student revision and missed-lecture recovery as the two highest-value use cases; defined and shipped **AI Video Lecture Generator** — pipeline extracts and structures teacher-recorded sessions into personalised async content.
- Contributed to **product roadmap** in collaboration with co-founders; maintained cross-functional alignment between product, engineering, and operations throughout both build cycles.

Software Development Intern

May 2025 – June 2025

ScoreMe Solutions

- Scoped and shipped a **two-stage PDF classification pipeline** for fintech document processing — rule-based fast path for deterministic inputs, **Random Forest fallback** for ambiguous cases; **automated ~80% of document volume**, reducing analyst review queue to exception-only cases.
- Defined **confidence threshold** as an explicit **human-in-the-loop gate**: outputs withheld from **credit scoring pipelines** and flagged for review; extended coverage to scanned PDFs via **Tesseract OCR**, doubling addressable input types.

Engineering Intern

June 2025 – July 2025

Bharat Electronics Limited

- Built a real-time stream processor for **ASTERIX CAT048** radar protocol per **EUROCONTROL** spec; decoded variable-length packet structures into structured aviation records; designed **concurrent ingestion** with **backpressure handling** to sustain throughput under high-frequency surveillance loads.

PROJECTS

CareerAutomated | Python, Groq/LLaMA, LangGraph, SQLite, Pandas, Playwright, IMAP/SMTP

2025

- Building an **autonomous AI recruiting platform** that orchestrates job discovery, candidate-job matching, resume tailoring, recruiter outreach, **ATS application execution**, CRM tracking, and inbox intelligence through a unified **multi-agent workflow**.
- Designed a **Generator-Critic architecture**: specialised agents handle company intelligence, project selection, and personalised email generation; a **Critic agent** validates quality, contextual relevance, formatting constraints, and placeholder safety before any outreach is sent — bad emails are blocked, not just flagged.
- Engineered **scalable data pipelines** using **SQLite**, **Pandas**, **IMAP**, and **SMTP** to maintain recruiter state, **deduplicate 10,000+ contacts**, synchronise Gmail Sent history, and automate personalised outreach at scale; **Playwright**-based browser automation executes ATS-specific application flows end-to-end.

YAAR — AI Behavioral Companion | React Native, LLM APIs, FastAPI, EAS Build

2025

- Shipped Android APK** (React Native + EAS Build) — chat UI, tap-based reaction probes, adaptive LLM generation, and shareable identity card export; **sub-2s streamed response latency** via FastAPI; validated end-to-end on device.
- Designed a **zero-input personalisation loop** (Open → Hook → React → Preference Update) for Gen Z in Tier 2/3 India; preference engine tracks **3 personality dimensions** from tap reactions, narrowing LLM output to identity-calibrated responses without retraining.
- Primary revenue via **contextual native brand integration** — behavioral state triggers in-character brand moments at peak-receptivity windows; freemium tier (50 reactions/day free vs. Rs.49/mo premium) as secondary.

Semantic Document Search — GDSC IIT Roorkee | Python, LangChain, BGE-M3, AstraDB, FastAPI, AWS EC2

2024

- Replaced keyword search on a **500+ document corpus** with **hybrid RAG** (dense **BGE-M3** + sparse BM25) — eliminated zero-result failure mode that caused query abandonment; users receive grounded direct answers instead of document lists.
- Offline **BGE-M3** embedding pipeline into **AstraDB** decoupled from **FastAPI serving on AWS EC2** — sub-second query response; corpus updates without downtime; prompt constraints enforce **response groundedness**.

TECHNICAL SKILLS

AI/ML: LangGraph, LangChain, Groq/LLaMA, MCP, Hybrid RAG, Multi-Agent Systems, Prompt Engineering, Generator-Critic Architecture, Human-in-the-Loop Design

Product: PRD Writing, MVP Scoping, Customer Discovery, Roadmapping, Prioritisation (RICE, Knapsack), A/B Testing, Funnel Analysis

Backend & Infra: Python, FastAPI, Docker, AWS EC2, Playwright, IMAP/SMTP, Stream Processing, REST APIs

Data & Analytics: SQL, SQLite, Pandas, Cohort Queries, OOP, DSA (graphs, trees, complexity)